16720-A S18 OCR using Neural Networks

Instructor: Kris Kitani

TAs: Leonid, Mohit, Arjun, Rawal, Aashi, Tanya

Due March 8, 2018 11:59 PM

Total Points: 100

Extra Credit Points: 20

Instructions

1. **Dropbox Link:** The data files in this homework are shared through the following dropbox link due to size restrictions. For any possible updates on the handout please check piazza.
   `https://www.dropbox.com/sh/ch8uwrgknz2ddlv/AAC1CMYMxC7H-NDh2Zoeitbva?dl=0`

2. **Integrity and collaboration:** Students are encouraged to work in groups but each student must submit their own work. If you work as a group, include the names of your collaborators in your write up. Code should **NOT** be shared or copied. Please **DO NOT** use external code unless permitted. Plagiarism is strongly prohibited and may lead to failure of this course.

3. **Start early!** This one is lengthy, training a network takes time and Q4.2 is tricky to implement. If you start late, you will be pressed for time while debugging. Also, speed read the entire assignment before you start so you get the gist of the assignment.

4. **Questions:** If you have any question, please look at piazza first. Other students may have encountered the same problem, and is solved already. If not, post your question on the discussion board. TAs will respond as soon as possible.

5. **Write-up:** Please note that we **DO NOT** accept handwritten scans for your write-up in this assignment. Please type your answers to theory questions and discussions for experiments electronically. Any word processor is allowed.

6. **Submission:** The submission is on Gradescope, **you will be submitting both your writeup and code zip file**. The zip file, `<andrew-id.zip>`, contains your MATLAB implementations (including helper functions), results for extra credit (optional). Do not handin the data files we distributed in the handout zip. **Note: You have to submit your writeup separately to Gradescope as `<andrew-id.pdf>`.**

   Your final upload should have the files arranged in this layout:

   `<AndrewID>.zip`

- `<AndrewID>/`
    - `matlab/`
        - `Backwards.m` (10 points)
        - `Classify.m` (2 points)
        - `ComputeAccuracyAndLoss.m` (3 points)
        - `extractImageText.m` (10 points)
        - `findLetters.m` (10 points)
        - `finetune36.m` (2 points)
        - `Forward.m` (5 points)
        - `InitializeNetwork.m` (5 points)
        - `Train.m` (5 points)
        - `testFindLetters.m` (5 points)
        - `testExtractImageText.m` (5 points)
        - `train26.m` (5 points)
        - `UpdateParameters.m` (5 points)
        - `define_autoencoder.m` (2 points)
        - `train_autoencoder.m` (4 points)
        - `eval_autoencoder.m` (4 points)
        - Any other code or helper functions you used
    - `ec/`(this is optional)
        - `checkGradient.m` (10 points)
        - `autoencoder_PCA.m` (10 points)
        - Any other extra credit code

7. TAs responsible for this assignment: Mohit Sharma (mohits1@andrew.cmu.edu) and Rawal Khirodkar (rkhirodk@andrew.cmu.edu).

Figure 1: Samples from NIST Special 19 dataset [1]

1 Overview

Deep learning has quickly become one of the most applied machine learning techniques in computer vision. Convolutional neural networks have been applied to many different computer vision problems such as image classification, recognition, and segmentation with great success. In this assignment, you will first implement a fully connected feed forward neural network for hand written character classification. Then in the second part, you will implement a system to locate characters in an image, which you can then classify with your deep network. The end result will be a system that, given an image of hand written text, will output the text contained in the image.

1.1 Basic Use

Here we will give a brief overview of the math for a single hidden layer feed forward network. For a more detailed look at the math and derivation, please see the class slides.

A fully-connected network $\mathbf{f}$, for classification, applies a series of linear and non-linear functions to an input data vector $\mathbf{x}$ of size $N \times 1$ to produce an output vector $\mathbf{f}(\mathbf{x})$ of size $C \times 1$, where each element i of the output vector represents the probability of $\mathbf{x}$ belonging to the class i. Since the data samples are of dimensionality N, this means the input layer has

Figure 2: Example of a single hidden layer network

N input units. To compute the value of the output units, we must first compute the values of all the hidden layers. The first hidden layer *pre-activation* $\mathbf{a}^{(1)}(\mathbf{x})$ is given by

$$\mathbf{a}^{(1)}(\mathbf{x}) = \mathbf{W}^{(1)}\mathbf{x} + \mathbf{b}^{(1)}$$

Then the *post-activation* values of the first hidden layer $\mathbf{h}^{(1)}(\mathbf{x})$ are computed by applying a non-linear activation function $\mathbf{g}$ to the *pre-activation* values

$$\mathbf{h}^{(1)}(\mathbf{x}) = \mathbf{g}(\mathbf{a}^{(1)}(\mathbf{x})) = \mathbf{g}(\mathbf{W}^{(1)}\mathbf{x} + \mathbf{b}^{(1)})$$

Subsequent hidden layer $(1 < t \leq T)$ pre- and post activations are given by:

$$\mathbf{a}^{(t)}(\mathbf{x}) = \mathbf{W}^{(t)}\mathbf{h}^{(t-1)} + \mathbf{b}^{(t)}$$
$$\mathbf{h}^{(t)}(\mathbf{x}) = \mathbf{g}(\mathbf{a}^{(t)}(\mathbf{x}))$$

The output layer *pre-activations* $\mathbf{a}^{(T)}(\mathbf{x})$ are computed in a similar way

$$\mathbf{a}^{(T)}(\mathbf{x}) = \mathbf{W}^{(T)}\mathbf{h}^{(T-1)}(\mathbf{x}) + \mathbf{b}^{(T)}$$

and finally the *post-activation* values of the output layer are computed with

$$\mathbf{f}(\mathbf{x}) = \mathbf{o}(\mathbf{a}^{(T)}(\mathbf{x})) = \mathbf{o}(\mathbf{W}^{(T)}\mathbf{h}^{(T-1)}(\mathbf{x}) + \mathbf{b}^{(T)})$$

where $\mathbf{o}$ is the output activation function. Please note the difference between $\mathbf{g}$ and $\mathbf{o}$!

For this assignment, we will be using the sigmoid activation function for the hidden layer, so:

$$\mathbf{g}(y) = \frac{1}{1 + \exp(-y)}$$

where when $\mathbf{g}$ is applied to a vector, it is applied element wise across the vector.

Since we are using this deep network for classification, a common output activation function to use is the softmax function. This will allow us to turn the real value, possibly negative values of $\mathbf{a}^{(T)}(\mathbf{x})$ into a set of probabilities (vector of positive numbers that sum to 1). Letting $\mathbf{x}_i$ denote the i^{th} element of the vector $\mathbf{x}$, the softmax function is defined as:

$$\mathbf{o}_i(\mathbf{y}) = \frac{\exp(\mathbf{y}_i)}{\sum_j \exp(\mathbf{y}_j)}$$

Q1.1.1 Theory [**2 points**] In training deep networks ReLU activation function is generally preferred to sigmoid. Why might this be the case?

Q1.1.2 Theory [**3 points**] All types of deep networks use non-linear activation functions for their hidden layers. Suppose we have a neural network with input dimension N and output dimension C and T hidden layers. Prove that if we have a linear activation function $\mathbf{g}$, then the number of hidden layers has no effect on the representation capability of the network (i.e., that the set of functions that can be represented by a T layer network is exactly the same as the set that can be represented by a $T' \neq T$ layer network).

2 Implement a Fully Connected Network

In this section, you will implement all of the functions needed to initialize, train, evaluate, and use the network. Throughout this assignment, the network will be represented by its parameters W and b, the weights and biases of the network.

2.1 Network Initialization

In order to use a deep network, we must first initialize the weights and biases in the network. This is typically done with a random initialization, or initializing the weights from some other training procedure (which we will see in Section 3). If you are interested in issues relating to initialization, [2] is a good paper to look at.

Q2.1.1 Theory [**2 points**]   Why is it not a good idea to initialize a network with all zeros? How about all ones, or some other constant value? (Hint: Consider what the gradients from backpropagation will look like.)

Q2.1.2 Code [**5 points**]   Implement the following function

```
[W, b] = InitializeNetwork(layers).
```

This function should take as input the sizes of the layers for your neural network. The `layers` parameters will be a vector of at least 3 integers. The first element denotes the size of the data layer N, the last element denotes the size of the output layer/number of classes C, and the intermediate elements denote the size of the hidden layers. You should return W and b which are both cell arrays containing `length(layers)-1` elements. W should contain at least two weight matrices of appropriate size, and b should contain at least two bias vectors.

Note: Use W or $b = cell(1,$num of layers$-1)$, we have provided trained network weights in this format.

Q2.1.3 Writeup [3 points] Describe the initialization you implemented in Q2.1.2 and any reasoning behind why you chose that strategy.

2.2 Forward Propagation

Section 1 has the math for forward propagation, we will implement it here. The loss function generally used for classification is the cross-entropy loss.

$$L_{\mathbf{f}}(\mathbf{D}) = - \sum_{(\mathbf{x},\mathbf{y}) \in \mathbf{D}} \log(\mathbf{y} \cdot \mathbf{f}(\mathbf{x}))$$

Here $\mathbf{D}$ is the full training dataset of data samples $\mathbf{x}$ ($N \times 1$ vectors, N = dimensionality of data) and labels $\mathbf{y}$ ($C \times 1$ one-hot vectors, C = number of classes).

Q2.2.1 Code [5 points]

$$[\text{out, act\_a, act\_h}] = \text{Forward(W, b, X)}$$

The function runs forward propagation on an input data sample X, using the network defined by the parameters W and b. Both W and b are cell arrays containing the network parameters as initialized by `InitializeNetwork(..)`.

X is a vector of dimensions $N \times 1$. This function should return the final softmax output of the network in the variable `out`, which is of size $C \times 1$. The function must also recturn the cell arrays `act_a` and `act_h`, which contain the pre- and post-activations of the network on this input sample.
Note: You might want to zero out any NaNs in this function using isnan().

Q2.2.2 Code [2 points]

$$[\text{outputs}] = \text{Classify(W, b, data)}$$

This function should accept the network parameters W and b, as well as a $D \times N$ matrix of data samples. `outputs` should contain the softmax output of the network for each of the data samples.

Q2.2.3 Code [3 points]

```
[accuracy, loss] = ComputeAccuracyAndLoss(W, b, data, labels)
```

This function should compute the accuracy on the network with respect to the provided `data` matrix of size $D \times N$ and `labels` matrix of size $D \times C$. You should also compute and return the average cross-entropy loss on the dataset.

2.3 Backwards Propagation

Gradient descent is an iterative optimisation algorithm, used to find the local optima. To find the local minima, we start at a point on the function and move in the direction of negative gradient (steepest descent) till some stopping criteria is met.

The update equation for a general weight $W_{ij}^{(t)}$ and bias $b_i^{(t)}$ is

$$W_{ij}^{(t)} = W_{ij}^{(t)} - \alpha * \frac{\partial L_{\mathbf{f}}}{\partial W_{ij}^{(t)}}(\mathbf{x}) \qquad b_i^{(t)} = b_i^{(t)} - \alpha * \frac{\partial L_{\mathbf{f}}}{\partial b_i^{(t)}}(\mathbf{x})$$

α is the learning rate. Please refer to the backpropagation slides for more details on how to derive the gradients. Note that here we are using softmax loss (which is different from the least square loss in the slides).

Q2.3.1 Code [10 points]

```
[grad_W, grad_b] = Backward(W, b, X, Y, act_h, act_a)
```

This function runs the back propogation algorithm on a single input example `X` and the ground truth vector `Y`. The function takes as input the network parameters `W` and `b` and the network pre- and post-activations `act_a` and `act_h` from calling `Forward` on `X`. The function should return the computed gradient updates for the network parameters. `grad_W` and `grad_b` are both cell arrays identical in shape to the kinds produced by `InitializeNetwork(..)`, but containing the gradient updates to make to each of the network parameters .

Q2.3.2 Code [2 points]

```
[W, b] = UpdateParameters(W, b, grad_W, grad_b, learning_rate)
```

Here we compute and return the updated network parameters `W` and `b`. The function is given the old parameters, the parameters gradients `grad_W` and `grad_b` returned by `Backward(..)`, and the supplied `learning_rate`.

2.4 Training Loop

When training a neural network using Gradient Descent, there are two options one can take for the updates, batch or stochastic updates. In batch gradient descent, we compute the gradient for all the examples in the training set and then average the gradients before updating the weights. In stochastic, we compute the gradient with one sample, update the weights with that gradient and carry on with the next sample. Each update is called a *iteration*, one complete pass through the entire dataset is called an *epoch*. In case of stochastic gradient descent, number of iterations in an epoch equals the number of training samples.

Q2.4.1 Theory [3 points] Give one pro and con for both stochastic and batch gradient descent.

Q2.4.2 Code [5 points]

```
[W, b] = Train(W, b, train_data, train_label, learning_rate)
```

The function trains the network for one epoch on `train_data` and ground-truth `train_label`. The function should return the updated network parameters.

2.5 Gradient Checker (Extra Credit)

Often, when implementing new layers in popular packages like Caffe or PyTorch you will have to implement the forward and backward pass for that layer. While the forward is fairly easy, a lot of things can go wrong in the backward pass. One common way to check if the backward pass is correct is by writing a gradient checker. A good explanation of a gradient checker is in the link: `http://ufldl.stanford.edu/tutorial/supervised/DebuggingGradientChecking/`

Q2.5.1 Extra Credit **[10 points]** Implement a script `checkGradient.m` that checks gradients of the loss with respect to a few random weights in each layer. Submit the code in the ec folder and report your results in the writeup. Specifically:

- For each layer,

  - randomly select a dimension in $\mathbf{W}$ ($\mathbf{b}$ is held constant)
  - add delta to that dimension, compute the loss. Now subtract delta and again compute the loss.
  - Use finite differences to estimate the gradient numerically. Compare this with gradient value of that dimension using your implementation of Backward( ).
  - repeat using a randomly selected dimension of $\mathbf{b}$ ($\mathbf{W}$ is held constant)
  - Add the $\mathbf{W}$ error and $\mathbf{b}$ error, this is your error for that layer

- Average out the error across layers and print it.

3 Training Models

We now have all the code required to train and use a deep network. In this section, you will train and evaluate different models which will then be used in Section 4 for parsing text in images.

3.1 From Scratch

Often times when you encounter new problems in deep learning, you will need to train a new network from scratch. This means that you will randomly initialize the weights in some way, and run multiple iterations of back propagation. The `Train.m` function you implemented in

Figure 3: Samples weights from a network with 50 hidden units trained for 30 epochs. Your weights will likely look very different, as you are training with more hidden units.

Q2.4.2 runs one epoch of training, passing over all data samples once. To fully train a deep network you typically must do many epochs.

Once a deep network is trained, one way to gain some insight into what it is doing is to visualize the learned weights. With convolutional neural networks as seen in class, each of the learned filters can be visualized as an image. This can also be done with a fully connected network, as you have implemented here. Since our input images are 32×32 images, unrolled into one 1024 dimensional vector that gets multiplied by $\mathbf{W}^{(1)}$, each row of $\mathbf{W}^{(1)}$ can be seen as a weight image. Reshaping each row into a 32×32 image can give us an idea of what types of images each unit in the hidden layer has a high response to.

We have provided you three data `.mat` files to use for this section. The training data in `nist26_train.mat` contains 299 samples for each of the 26 upper-case letters of the alphabet. This is the set you should use for training your network. The cross-validation set in `nist26_valid.mat` contains 100 samples from each class, and should be used in the training loop to see how the network is performing on data that it is not training on. This will help to spot over fitting. Finally, the test data in `nist26_test.mat` contains another 100 samples per class, and should be used for the final evaluation on your best model to see how well it will generalize to new unseen data.

Q3.1.1 Code [3 points] Use the provided `train26.m` to train a network from scratch. Use a single hidden layer with 400 hidden units, and train for at least 30 epochs. **Modify** the script to plot generate two plots: one showing the accuracy on both the training and validation set over the epochs, and the other showing the cross-entropy loss averaged over the data. The x-axis should represent the epoch number, while the y-axis represents the accuracy or loss. With these settings, you should see an accuracy on the validation set of at least 75%.

Figure 4: Learned filters from AlexNet trained on the ImageNet dataset. AlexNet is a convolutional neural network who's architecture and weights are often used as the basis for new networks.

Q3.1.2 Writeup [2 points]   Use your modified training script to train two networks, one with learning rate 0.01, and another with learning rate 0.001. Include all 4 plots in your writeup. Comment on how the learning rates affect the training, and report the final accuracy of the best network on the test set.

Q3.1.3 Writeup [3 points]   Using the best network from the previous question, report the accuracy and average cross-entropy loss on the test set, and visualize the first layer weights that your network learned (using `reshape` and `montage`). Compare these to the network weights immediately after initialization. Include both visualizations in your writeup. Comment on the learned weights. Do you notice any patterns?

Q3.1.4 Writeup [2 points]   Visualize the confusion matrix for your best model as a 26×26 image (upscale the image so we can actually see it). Comment on the top two pairs of classes that are most commonly confused.

3.2   Fine Tuning

When training from scratch, a lot of epochs are often needed to learn anything meaningful. One way to avoid this is to instead initialize the weights more intelligently. Various strategies have been used for initializing neural networks, such as unsupervised pretraining with Auto-Encoders or Restricted Boltzmann Machines.

However thanks to the explosion in popularity of deep learning for computer vision, it is often possible to also initialize a network with weights from another deep network that was trained for a different purpose. This is because, whether we are doing image classification, segmentation, recognition etc..., most real images share common properties. Simply copying the weights from the other network to yours gives your network a head start, so your network does not need to learn these common weights from scratch all over again. This is also referred to as fine tuning.

We have trained a network for recognizing capital letters using 800 hidden units, and trained it for 60 epochs. The network parameters are stored in `nist26_model_60iters.mat`. Using these weights, you will initialize and fine-tune a network for recognizing both capital letters as well as numeric digits.

Q3.2.1 Code/Writeup [2 points]   Make a copy of `train26.m` and name it `finetune36.m`. Modify this script to load the data from `nist36_*.mat`, and train a network to classify both written letters and numbers. Finetune (train) this network for 5 epochs with learning rate 0.01, and include plots of the accuracy and average cross-entropy loss in your writeup.

Q3.2.2 Writeup [2 points]   Once again, visualize the network's first layer weights before and after training. Comment on the differences you see. Also report the network's accuracy and average loss on the test set.

Q3.2.3 Writeup [1 points]   Visualize the confusion matrix for your best model as a 36×36 image (upscale the image so we can actually see it). Comment on the top two pairs of classes that are most commonly confused. How has introducing more classes affected the network?

4   Extract Text from Images

Now that you have a network that can recognize handwritten letters with reasonable accuracy, you can now use it to parse text in an image. Given an image with some text on it, our goal is to have a function that returns the actual text in the image. However, since your neural network expects a a binary image with a single character, you will need to process the input image to extract each character. There are various approaches that can be done so feel free to use any strategy you like.
   Here we outline one possible method:

1. Process the image (blur, threshold, etc...) to classify all pixels as being part of a character or background.

2. Find connected groups of character pixels (see `bwconncomp`). Place a bounding box around each connected component.

3. Group the letters based on which line of the text they are a part of, and sort each group so that the letters are in the order they appear on the page.

Figure 5: Sample image with handwritten characters annotated with boxes around each character.

4. Take each bounding box one at a time and resize it to 32×32, classify it with your network, and report the characters in order (inserting spaces when it makes sense).

Since the network you trained likely does not have perfect accuracy, you can expect there to be some errors in your final text parsing. Whichever method you choose to implement for the character detection, you should be able to place a box on most of there characters in the image. We have provided you with `01_list.jpg`, `02_letters.jpg`, `03_haiku.jpg` and `04_deep.jpg` to test your implementation on.

Q4.1 Theory [2 points] The method outlined above is pretty simplistic, and makes several assumptions. What are two big assumptions that the sample method makes. In your writeup, include two example images where you expect the character detection to fail (either miss valid letters, or respond to non-letters).

Q4.2 Code [10 points]

$$[\texttt{lines, bw}] = \texttt{findLetters(im)}$$

Given an RGB image, this function should return a cell array `lines` containing all of the located handwritten characters in the image, as well as a binary black-and-white version of the image `im`. The `letters` cell array should contain one entry for each line of text that appears in the image. Each entry should be a matrix of size $L_i \times 4$, where L_i represents the number of characters in the i^{th} line of text. Each row of the matrix should contain `[x1, y1, x2, y2]` the positions of the top-left and bottom-right corners of the box. The processed image `bw` should be only black and white (containing pixel values 0.0 and 1.0), with the characters in black and the rest of the image in white. Also include a test script `testFindLetters.m` that runs your function on all of the provided images (and any others you decided to include), and displays a figure showing the bounding boxes around each letter.

Q4.3 Writeup [3 points] Run `findLetters(..)` on all of the provided sample images in `images/`. Plot all of the located boxes on top of the image to show the accuracy of your `findLetters(..)` function. Include all the result images in your writeup.

Q4.4 Code [10 points]

$$[\text{text}] = \text{extractImageText(fname)}.$$

This function should accept an image path `fname`, and return the text contained in the image. This function will load the image, find the character locations, classify each one with the network you trained in **Q3.2.1**, and return the text contained in the image. Also include a test script `testExtractImageText.m` that runs your function on all of the provided images (and any others you decided to include), and outputs the retrieved text for each image.

Q4.5 Writeup [5 points] Run your `extractImageText(..)` on all of the provided sample images in `images/`. Include the extracted text in your writeup.

5   Image Compression with Autoencoders

Note: You need to install the newest MATLAB 2017b with the (new) Neural Network Toolbox. You will not be able to do this section with MATLAB 2017a.

Deep learning is not only useful for supervised learning such as classification; they also perform well in *unsupervised learning* scenarios. In this section, you will explore an important type of network architecture in deep learning research called *autoencoders*. An autoencoder is a neural network that is trained to attempt to copy its input to its output, but it usually allows copying only approximately. This is typically achieved by restricting the number of hidden nodes inside the autoencoder; in other words, the autoencoder would be forced to learn to *represent* data with this limited number of hidden nodes. This is a very useful way of learning compressed representations. In this section, we will continue using the NIST36 dataset you have from the previous questions.

Additionally, this section will help you understand how matlab implements its deep learning toolbox. Most other libraries including TensorFlow, PyTorch, Caffee *etc.* also have very similar setups and API's. However, we won't be covering them here.

5.1   Building the Autoencoder [2 points]

Take a look at the provided code `define_autoencoder.m`. We are going to build a simple autoencoder with the serial architecture.

- 2D convolution: 4 filters of size 4×4, padding=[1,1], stride=[2,2]

- 2D convolution: 8 filters of size 4×4, padding=[1,1], stride=[2,2]

- 2D convolution: 64 filters of size 8×8, no padding, stride=[1,1]

- 2D transposed convolution: 8 filters of size 8×8, no cropping, stride=[1,1]

- 2D transposed convolution: 4 filters of size 4×4, cropping=[1,1], stride=[2,2]

- 2D transposed convolution: 1 filter of size 4×4, cropping=[1,1], stride=[2,2]

Note that the output feature of the last 2D convolution will be of size $1 \times 1 \times 64$.

Notice that this roughly follows a symmetric structure. The output dimensions will be exactly the same as the input dimensions. We will use our input images to supervise the output of the network and optimize with pixel-wise L_2 loss; we do not need any labels. *Transposed convolution* here (often misinterpreted as deconvolution) is a similar operation to convolution, but the filters are for increasing spatial resolutions rather than reducing them. They are useful operations for generative models that creates data from hidden representations.

Q5.1 Writeup/Code [2 points] There is something wrong with `define_autoencoder.m`. What is it? Correct the code, and explain how you did it.

5.2 Training the Autoencoder [4 points]

After defining the autoencoder, let's implement `train_autoencoder.m`. You will need to (1) load the training data and (2) reshape/adjust the dimensions appropriately before it can be read into the network, and (3) finally, complete the unfinished call to the built-in function `trainNetwork()`.

Q5.2.1 Writeup/Code [2 points] Using the provided default settings, train the network for 3 epochs. What do you observe in the plotted training loss curve as it progresses?

There are various factors that can impact the training of neural networks. Two of the most important ones are the *learning rate* and the *weight initialization*. Adjust these two hyper-parameters to make the training loss decrease faster. As a reference, a good training setting in this problem would make the training mini-batch loss decrease to around 4 after the 3rd epoch.

Q5.2.2 Writeup [2 points] What adjustments did you make to improve the training performance? Include a snapshot of the plotted training progress in your report.

5.3 Evaluating the Autoencoder [4 points]

Now lets evaluate how well the autoencoder has been trained. In `eval_autoencoder.m`, load the test data and perform the same data preprocessing as you did with the training data. Complete the unfinished call to the built-in function `predict()`. The output of `predict()` should give you the reconstructed version of the test images. Remember that these reconstructions can be represented with only 64 numbers instead of the original $32 \times 32 = 1024$.

Q5.3.1 Writeup/Code [2 points] Select 5 classes from the total 36 classes in your dataset and for each selected class include in your report 2 test images and their reconstruction. You may use test labels to help you find the corresponding classes. The function `subplot` will

also be useful. Since dimensionality reduction with autoencoders belong to the class of lossy compression, the exact original data will not be recovered. What differences do you observe that exist in the reconstructed test images, compared to the original ones?

Q5.3.2 Writeup [2 points] Lets evaluate the reconstruction quality using Peak Signal-to-noise Ratio (PSNR). PSNR is defined as

$$\text{PSNR} = 20 \times \log_{10}(\text{MAX}_I) - 10 \times \log_{10}(\text{MSE}) \tag{1}$$

where MAX_I is the maximum possible pixel value of the image, and MSE (mean squared error) is computed across all pixels. You may use the built-in function `psnr()` for convenience. Report the average PSNR you get from the autoencoder across all test images.

5.4 Comparing against PCA [10 points] - Extra-Credit

As a baseline for comparison, we will use one of the most popular methods for data dimensionality reduction - Principle Component Analysis (PCA). PCA allows one to find the best low-rank approximation of the data by keeping only a specified number of principle components. To perform PCA, we will use a factorization method called Singular Value Decomposition (SVD).

For further reading, please refer to

- https://en.wikipedia.org/wiki/Principal_component_analysis

- https://en.wikipedia.org/wiki/Singular-value_decomposition

Run `[U,S,V] = svd(train data)`. One of the matrices will be an orthonormal matrix that indicates the components of your data, sorted by their importances. Extract the first 64 principle components and form a projection matrix; you will need to figure out how to do these from the `U,S,V` matrices. All code that you write for this question should be added in `autoencoder_PCA.m`.

Q5.4.1 Writeup [2 points] What is the size of your projection matrix? What is its rank? This projection matrix was "trained" from our training data. Now lets "test" it on our test data. Use the projection matrix on test data to obtain the reconstructed test images. Note that these reconstructions can also be represented with only 64 numbers.

Q5.4.2 Writeup [2 points] Use the classes you selected in `Q5.3.1`, and for each of these 5 classes, include in your report 2 test images and their reconstruction. You may use test labels to help you find the corresponding classes. What differences do you observe that exist in the reconstructed test images, compared to the original ones? How do they compare to the ones reconstructed from your autoencoder?

Q5.4.3 Writeup [2 points] Report the average PSNR you get from PCA. Is it better than your autoencoder? Why?

Q5.4.4 Writeup [2 points]   Train another autoencoder with your current settings, but *without* the corrections you made in Q5.1. Include the same 2 test images and their reconstruction in your report, and report the average PSNR you get from the autoencoder and from PCA. What similarity/difference do you observe? Why is there such similarity/difference?

Q5.4.5 Writeup [2 points]   Count the number of learned parameters for both your autoencoder and the PCA model. How many parameters do the respective approaches learn? Why is there such a difference in terms of performance?

References

[1] P. J. Grother. Nist special database 19   handprinted forms and characters database. https://www.nist.gov/srd/nist-special-database-19, 1995.

[2] Xavier Glorot and Yoshua Bengio. Understanding the difficulty of training deep feedforward neural networks.